

Logistic Regression - Supervised Learning - Classification

1. Data preprocessing

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv('data.csv')
df.head()
```

```
Out[2]:
```

	-6	592	0
0	-5	807	0
1	-5	-344	0
2	-5	-126	0
3	-5	243	0
4	-5	185	0

```
In [3]: df.shape
```

```
Out[3]: (199, 3)
```

```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 199 entries, 0 to 198
Data columns (total 3 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    -6      199 non-null    int64
 1   592      199 non-null    int64
 2    0      199 non-null    int64
dtypes: int64(3)
memory usage: 4.8 KB

```

In [5]: `df.dtypes`

```

Out[5]: -6      int64
        592      int64
        0      int64
        dtype: object

```

In [6]: `df.describe()`

```

Out[6]:
```

	-6	592	0
count	199.000000	199.000000	199.000000
mean	0.964824	325.477387	0.582915
std	3.636761	286.265167	0.494321
min	-5.000000	-445.000000	0.000000
25%	-2.000000	152.500000	0.000000
50%	1.000000	349.000000	1.000000
75%	4.000000	538.000000	1.000000
max	8.000000	1056.000000	1.000000

In [7]: `df.isnull().sum()`

```
Out[7]: -6      0
        592     0
        0      0
        dtype: int64
```

```
In [8]: df[df.duplicated()]
```

```
Out[8]:
```

	-6	592	0
40	-3	421	1
169	5	454	1

```
In [9]: df.shape
```

```
Out[9]: (199, 3)
```

```
In [10]: df.drop_duplicates(inplace=True)
```

```
In [11]: df.shape
```

```
Out[11]: (197, 3)
```

```
In [12]: df.rename(columns={'-6': 'Feature1', '592': 'Feature2', '0': 'Label'}, inplace=True)
```

```
In [13]: df.head()
```

```
Out[13]:
```

	Feature1	Feature2	Label
0	-5	807	0
1	-5	-344	0
2	-5	-126	0
3	-5	243	0
4	-5	185	0

```
In [42]: df.to_csv('updated_df.csv')
```

```
In [15]: df.nunique()
```

```
Out[15]: Feature1      14  
Feature2     180  
Label         2  
dtype: int64
```

Observations 💡

- 3 Columns available only
- No Null Values
- 2 Duplicate values found
- Column Names changed accordingly

```
In [16]: for i in df.columns:  
    print(f'{i} Colum Unique Values :',end='')  
    print(df[i].unique())  
    print('*'*20)
```

```

Feature1 Colum Unique Values :[-5 -4 -3 -2 -1 0 1 2 3 4 5 6 7 8]
*****
Feature2 Colum Unique Values :[ 807 -344 -126 243 185 -89 -7 -90 161 35 611 -45 558 -110
168 349 91 674 697 12 426 7 -65 634 342 -39 314 679
627 272 358 -146 326 274 421 -237 465 707 104 622 408 454
93 683 -181 300 389 458 157 831 33 719 671 306 743 515
166 202 356 333 -68 355 416 240 270 -445 292 267 203 417
655 263 691 778 135 761 212 1056 490 66 575 623 -140 584
482 630 164 646 96 388 -206 331 95 28 -128 -66 -4 826
404 64 607 148 688 649 -187 -185 395 200 768 -338 282 668
222 182 285 735 452 619 400 399 414 579 319 57 352 658
-97 382 -298 378 276 740 564 269 -226 572 245 650 537 654
596 428 364 466 65 100 571 504 178 193 485 250 517 442
-281 339 653 110 -49 641 460 206 -322 219 336 748 249 287
528 469 273 73 481 390 80 375 959 539 727 75]
*****
Label Colum Unique Values :[0 1]
*****

```

2. Exploratory Data Analysis

```
In [17]: import seaborn as sns
```

```
In [18]: df.head()
```

```
Out[18]:
```


	Feature1	Feature2	Label
0	-5	807	0
1	-5	-344	0
2	-5	-126	0
3	-5	243	0
4	-5	185	0

```
In [19]: plt.figure(figsize=(15,4))
plt.suptitle('Histograms for Features')
```

```
plt.subplot(1,2,1)
sns.histplot(data = df,x = df['Feature1'],color='r')
plt.xlabel('Feature 1')
plt.ylabel('Frequency')
plt.title(f'Histogram for Feature 1')

plt.subplot(1,2,2)
sns.histplot(data = df,x = df['Feature2'])
plt.xlabel('Feature 2')
plt.ylabel('Frequency')
plt.title(f'Histogram for Feature 2')

plt.show()
```



- Observations - Feature2 Data is distributed normally


```
In [20]: plt.figure(figsize=(15,4))
plt.suptitle('Outliers for Features')

plt.subplot(1,2,1)
sns.boxplot(data = df,x = df['Feature1'],color='r')
```

```
plt.xlabel('Feature 1')
plt.ylabel('Frequency')
plt.title(f'Boxplot for Feature 1')

plt.subplot(1,2,2)
sns.boxplot(data = df,x = df['Feature2'],orient='v')
plt.xlabel('Feature 2')
plt.ylabel('Frequency')
plt.title(f'Boxplot for Feature 2')

plt.show()
```



```
In [21]: np.percentile(df['Feature1'],0.25),np.percentile(df['Feature1'],0.75),df['Feature1'].median()
```

```
Out[21]: (-5.0, -5.0, 1.0)
```

```
In [22]: np.percentile(df['Feature2'],0.25),np.percentile(df['Feature2'],0.75),df['Feature2'].median()
```

```
Out[22]: (-395.51, -341.18, 342.0)
```

```
In [23]: outlier = df[df['Feature2'] < np.percentile(df['Feature2'], 0.25)]
outlier
```


```
Out[23]:
```

	Feature1	Feature2	Label
--	----------	----------	-------

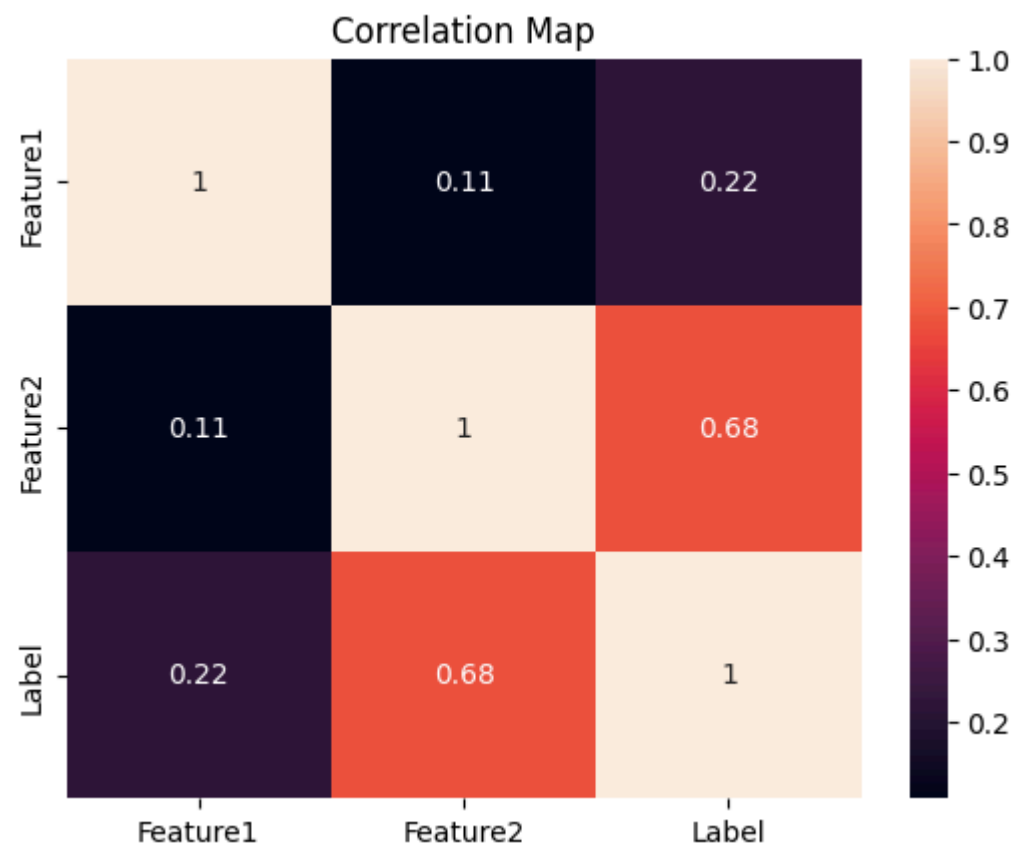
68	-1	-445	0
----	----	------	---

- Observations - Feature2 has outliers

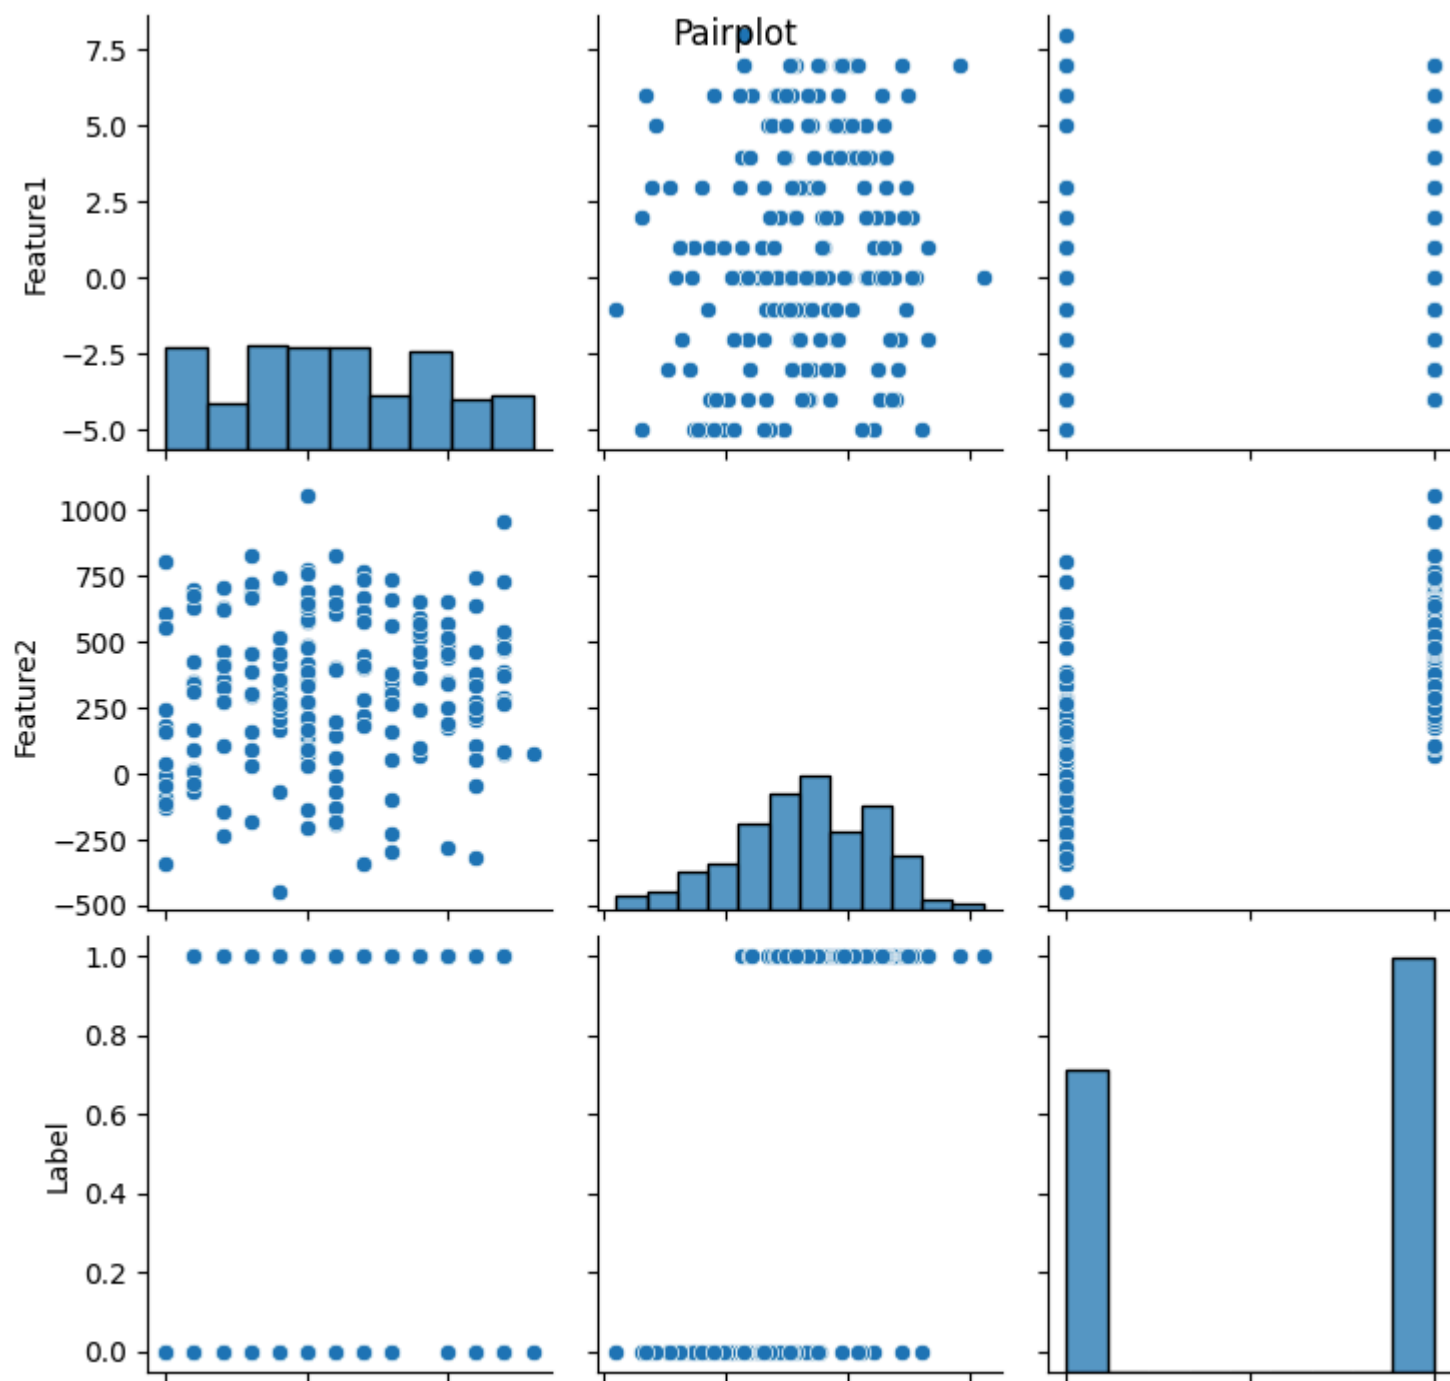
```
In [24]: sns.scatterplot(data=df, x='Feature1', y='Feature2', hue='Label')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.title('Scatterplot for Features')
plt.grid()
plt.show()
```

```
In [25]: sns.heatmap(df.corr(),annot=True,cmap='rocket')  
plt.title('Correlation Map')  
plt.show()
```



```
In [26]: sns.pairplot(df)
plt.suptitle('Pairplot')
plt.show()
```



−5 0 5 −500 0 500 1000 0.0 0.5 1.0
Feature1 Feature2 Label

```
In [27]: df.shape
```

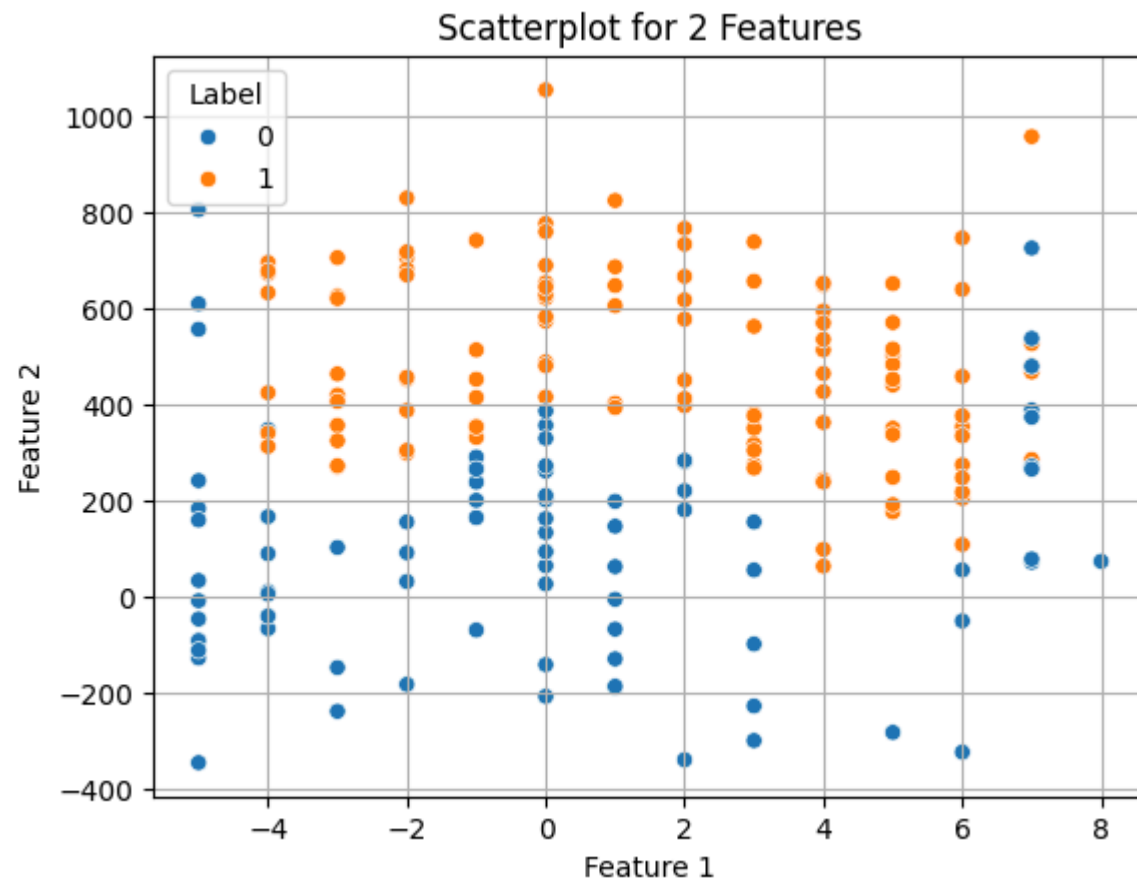
```
Out[27]: (197, 3)
```

```
In [28]: df.drop(outlier.index,inplace=True)
```

```
In [29]: df.shape
```

```
Out[29]: (196, 3)
```

```
In [30]: sns.scatterplot(data = df, x='Feature1',y='Feature2', hue="Label")  
plt.grid()  
plt.xlabel('Feature 1')  
plt.ylabel('Feature 2')  
plt.title('Scatterplot for 2 Features')  
plt.show()
```



```
In [32]: df.head()
```

```
Out[32]:
```

	Feature1	Feature2	Label
0	-5	807	0
1	-5	-344	0
2	-5	-126	0
3	-5	243	0
4	-5	185	0

```
In [33]: df.isnull().sum()
```

```
Out[33]: Feature1    0  
Feature2    0  
Label      0  
dtype: int64
```

```
In [34]: df.shape
```

```
Out[34]: (196, 3)
```

```
In [ ]: # Libraries  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import accuracy_score
```

A) Logistic Regression

```
In [75]: # Splitting  
X_train, X_test, y_train, y_test = train_test_split(df[['Feature1', 'Feature2']], df['Label'], test_size=0.2, random_state=42)  
  
# Standarization  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

```

# Initializing model
model_sklearn = LogisticRegression()

# Training
model_sklearn.fit(X_train_scaled, y_train)

# Decision boundary + plotting
def plot_decision_boundary_sklearn_scaled(X, y, model, title):
    h = .02 # Step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)


    plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.8)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Paired)
    plt.title(title)
    plt.xlabel('Feature1 ')
    plt.ylabel('Feature2 ')
    plt.show()

# For Trained Data
plot_decision_boundary_sklearn_scaled(X_train_scaled, y_train, model_sklearn, 'Decision Boundary')

# Evaluation
X_test_scaled = scaler.transform(X_test)
y_pred_sklearn = model_sklearn.predict(X_test_scaled)
accuracy_sklearn = accuracy_score(y_test, y_pred_sklearn)

# Metric
print(f"Model Accuracy : {accuracy_sklearn:.2%}")

```



Model Accuracy : 85.00%

B) Implementation of Logistic Regression from Scratch without scikit-learn

```
In [58]: def sigmoid(z):  
         return 1 / (1 + np.exp(-z))  
  
         def compute_cost(X, y, theta):  
             m = len(y)  
             h = sigmoid(X @ theta)  
             cost = (-1/m) * np.sum(y * np.log(h) + (1 - y) * np.log(1 - h))  
             return cost
```



```

def gradient_descent(X, y, theta, alpha, iterations):
    m = len(y)
    costs = []

    for i in range(iterations):
        h = sigmoid(X @ theta)
        gradient = (1/m) * X.T @ (h - y)
        theta -= alpha * gradient
        cost = compute_cost(X, y, theta)
        costs.append(cost)

    return theta, costs

X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

# Scaling
scaler = StandardScaler()
X_train_scaled[['Feature1', 'Feature2']] = scaler.fit_transform(X_train[['Feature1', 'Feature2']])
X_test_scaled[['Feature1', 'Feature2']] = scaler.transform(X_test[['Feature1', 'Feature2']])

# Bias
X_train_b = np.c_[np.ones((X_train_scaled.shape[0], 1)), X_train_scaled]
X_test_b = np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled]

initial_theta = np.zeros(X_train_b.shape[1])

# Parameters
learning_rate = 0.01
num_iterations = 1000

# Training
trained_theta, costs = gradient_descent(X_train_b, y_train, initial_theta, learning_rate, num_iterations)

# Cost Function For Every Iterations
plt.plot(range(1, num_iterations + 1), costs)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost Function during Training - (without scikit-learn)')
plt.show()

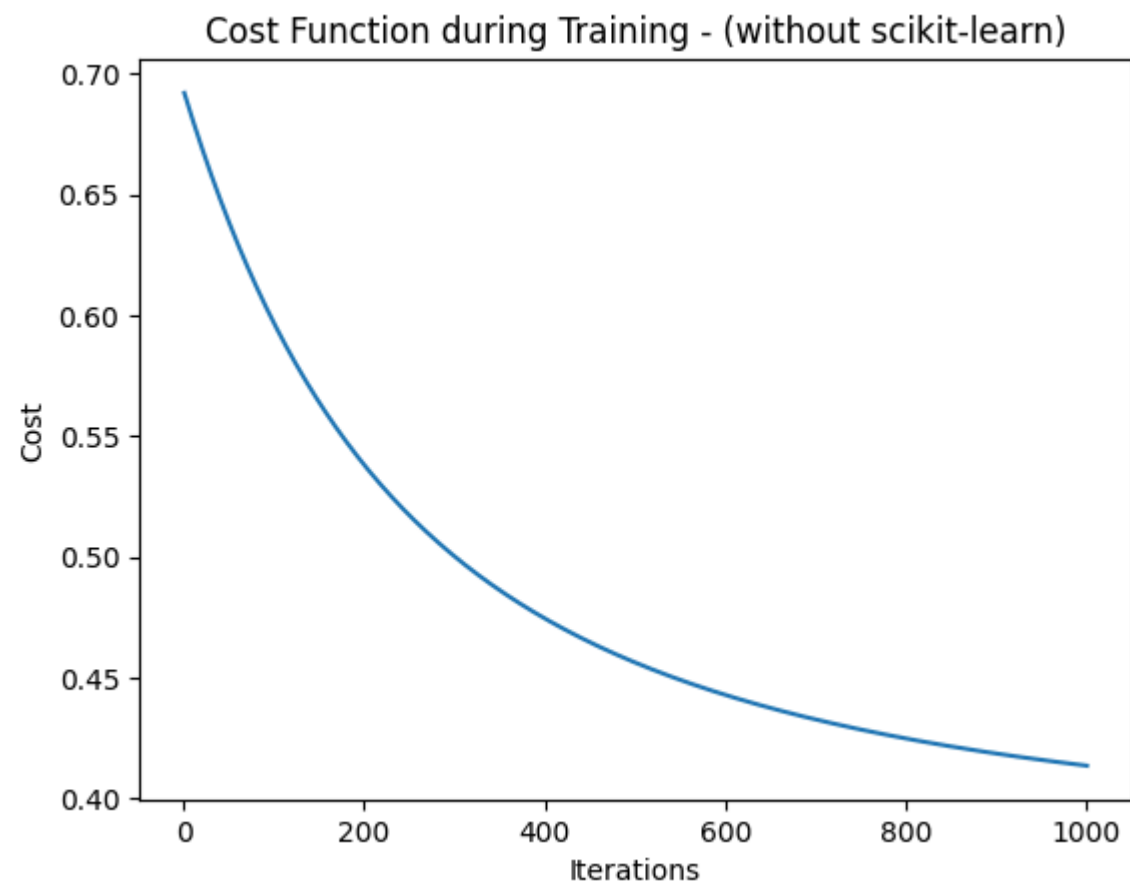
```

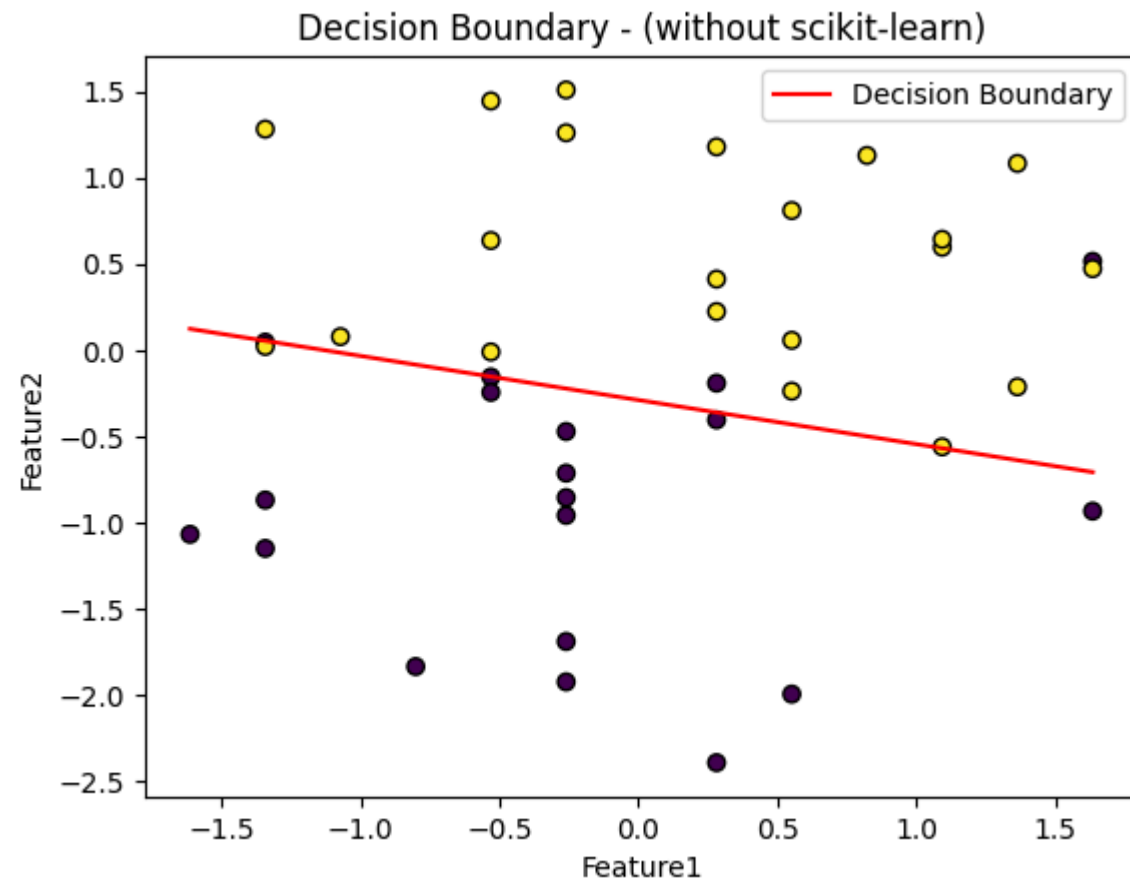
```

# Decision Boundary
x_values = np.linspace(X_test_scaled['Feature1'].min(), X_test_scaled['Feature1'].max(), 100)
plt.scatter(X_test_scaled['Feature1'], X_test_scaled['Feature2'], c=y_test, cmap='viridis', edgecolors='k', marker='o')
y_values = -(trained_theta[0] + trained_theta[1] * x_values) / trained_theta[2]
plt.plot(x_values, y_values, label='Decision Boundary', color='red')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.title('Decision Boundary - (without scikit-learn)')
plt.legend()
plt.show()

# Prediction & Accuracy
predictions = sigmoid(X_test_b @ trained_theta)
predictions = (predictions >= 0.5).astype(int)
accuracy = np.mean(predictions == y_test)
print(f"Model Accuracy: {accuracy*100:.2f}%")

```





Model Accuracy: 92.50%

C) Implementation of Logistic Regression + Feature Engineering

```
In [61]: df = pd.read_csv('updated_df.csv')
```

```
In [62]: df.head()
```

Out[62]:

	Unnamed: 0	Feature1	Feature2	Label
0	0	-5	807	0
1	1	-5	-344	0
2	2	-5	-126	0
3	3	-5	243	0
4	4	-5	185	0

In [63]: `df.drop(columns='Unnamed: 0',inplace=True)`

In [64]: `df.shape`

Out[64]: (196, 3)

In [65]: `df.isnull().sum()`

Out[65]:

```
Feature1    0
Feature2    0
Label       0
dtype: int64
```

In [73]: `from sklearn.preprocessing import PolynomialFeatures # for Increasing Degree`

```
poly = PolynomialFeatures(degree=2)
X_poly = poly.fit_transform(df[['Feature1', 'Feature2']])
poly_feature_names = [f'poly_{i}' for i in range(X_poly.shape[1])]
df_poly = pd.DataFrame(X_poly, columns=poly_feature_names)
df_poly['Label'] = df['Label']

# Splitting
X_train_poly, X_test_poly, y_train_poly, y_test_poly = train_test_split(
    df_poly.iloc[:, :-1], df_poly['Label'], test_size=0.2, random_state=42)

# Standardization
scaler_poly = StandardScaler()
X_train_scaled_poly = scaler_poly.fit_transform(X_train_poly)
```

```

X_test_scaled_poly = scaler_poly.transform(X_test_poly)

# Initializing model
model_poly_sklearn = LogisticRegression()

# Training
model_poly_sklearn.fit(X_train_scaled_poly, y_train_poly)

# Decision boundary
def plot_decision_boundary_sklearn_poly(X, y, model, title):
    h = .02
    x_min, x_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    y_min, y_max = X[:, 2].min() - 1, X[:, 2].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    meshgrid_poly = poly.transform(np.c_[xx.ravel(), yy.ravel()])

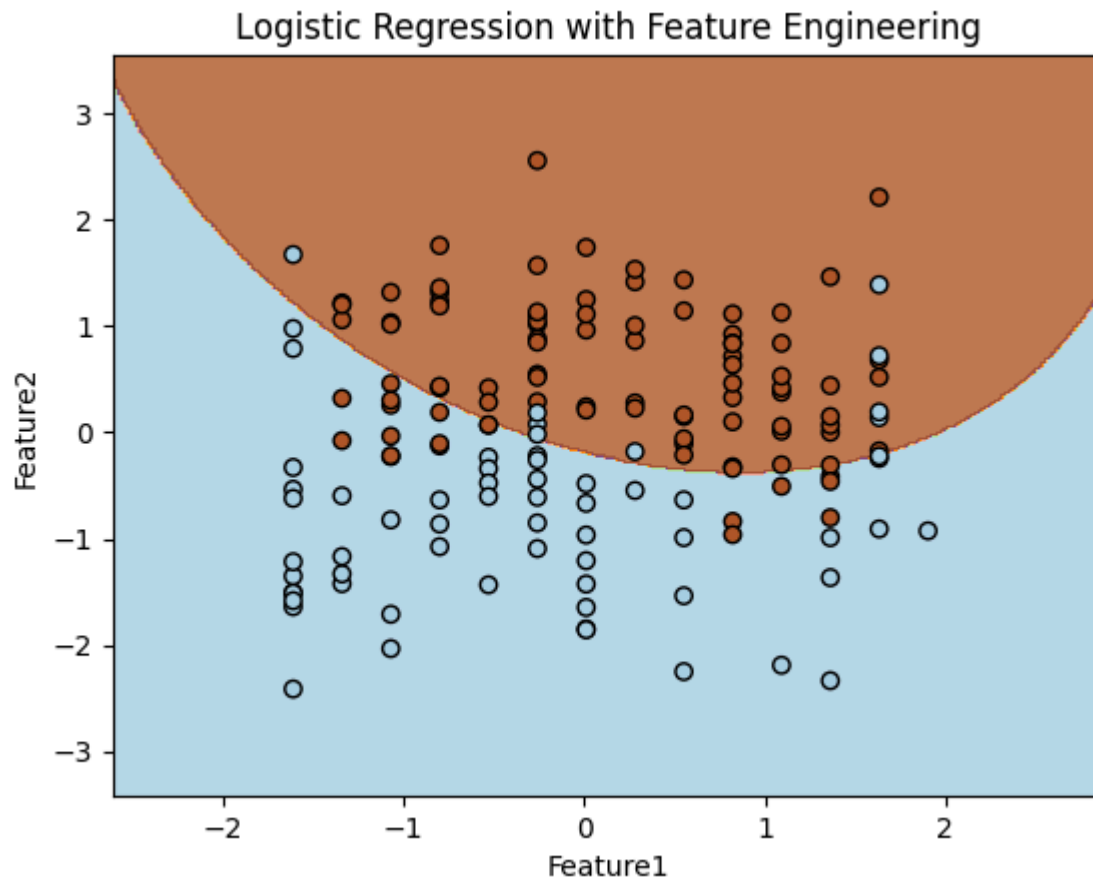
    Z = model.predict(meshgrid_poly)
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.8)
    plt.scatter(X[:, 1], X[:, 2], c=y, edgecolors='k', cmap=plt.cm.Paired)
    plt.title(title)
    plt.xlabel('Feature1')
    plt.ylabel('Feature2')
    plt.show()

plot_decision_boundary_sklearn_poly(X_train_scaled_poly, y_train_poly, model_poly_sklearn, 'Logistic Regression with Feature E

# Evaluation
y_pred_poly_sklearn = model_poly_sklearn.predict(X_test_scaled_poly)
accuracy_poly_sklearn = accuracy_score(y_test_poly, y_pred_poly_sklearn)
print(f"Model Accuracy with Feature Engineering: {accuracy_poly_sklearn:.2%}")

```



Model Accuracy with Feature Engineering: 82.50%

```
In [72]: X_train_poly.shape, X_test_poly.shape, y_train_poly.shape, y_test_poly.shape
```

```
Out[72]: ((156, 6), (40, 6), (156,), (40,))
```

D)Implementation of Logistic Regression from Scratch with Feature Engineering

```
In [71]: def sigmoid(z):  
         return 1 / (1 + np.exp(-z))  
  
         def compute_cost(X, y, theta):
```

```

    m = len(y)
    h = sigmoid(X @ theta)
    cost = (-1/m) * np.sum(y * np.log(h) + (1 - y) * np.log(1 - h))
    return cost

def gradient_descent(X, y, theta, alpha, iterations):
    m = len(y)
    costs = []

    for i in range(iterations):
        h = sigmoid(X @ theta)
        gradient = (1/m) * X.T @ (h - y)
        theta -= alpha * gradient
        cost = compute_cost(X, y, theta)
        costs.append(cost)

    return theta, costs

X_train_poly_scaled = X_train_poly.copy()
X_test_poly_scaled = X_test_poly.copy()

# Scaling
scaler_poly = StandardScaler()
X_train_poly_scaled = scaler_poly.fit_transform(X_train_poly)
X_test_poly_scaled = scaler_poly.transform(X_test_poly)

# Bias
X_train_poly_b = np.c_[np.ones((X_train_poly_scaled.shape[0], 1)), X_train_poly_scaled]
X_test_poly_b = np.c_[np.ones((X_test_poly_scaled.shape[0], 1)), X_test_poly_scaled]

initial_theta_poly = np.zeros(X_train_poly_b.shape[1])

# Parameters
learning_rate_poly = 0.01
num_iterations_poly = 1000

# Training
trained_theta_poly, costs_poly = gradient_descent(X_train_poly_b, y_train_poly, initial_theta_poly, learning_rate_poly, num_it

# Cost Function For Every Iterations
plt.plot(range(1, num_iterations_poly + 1), costs_poly)

```



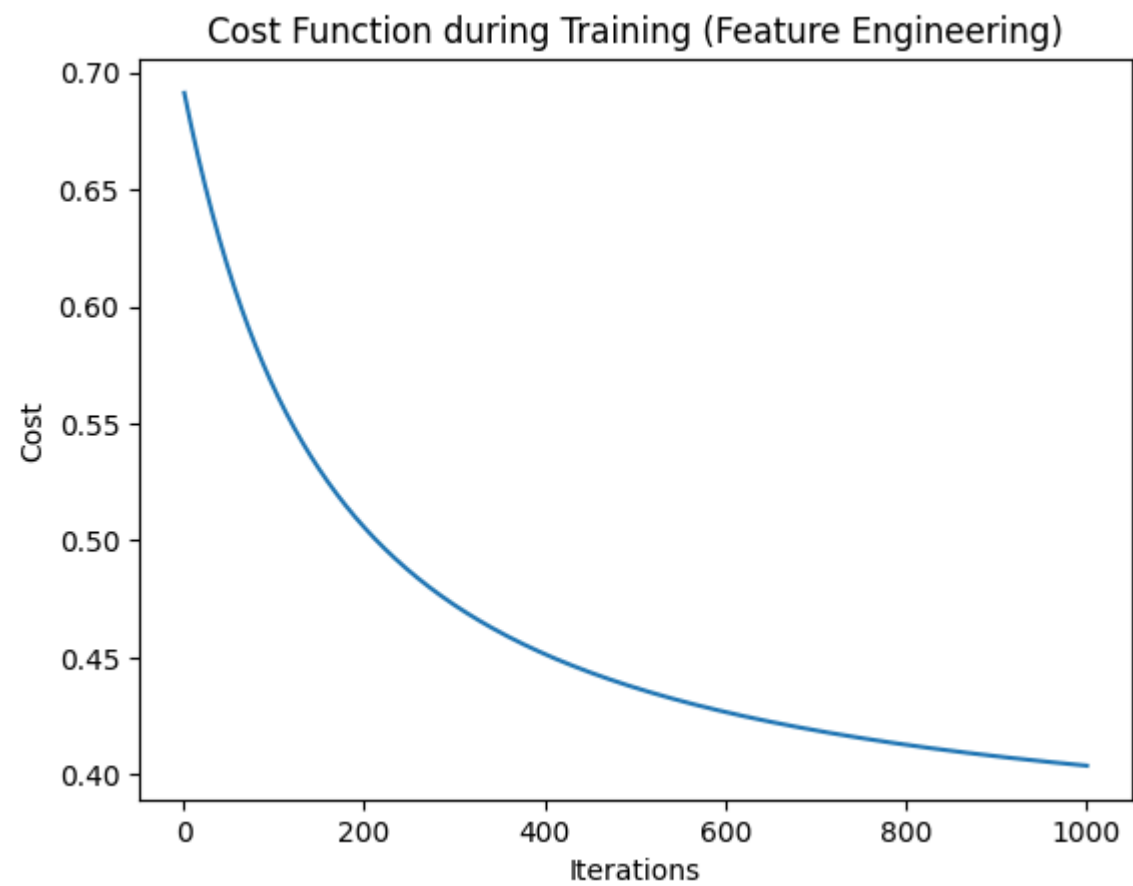
```

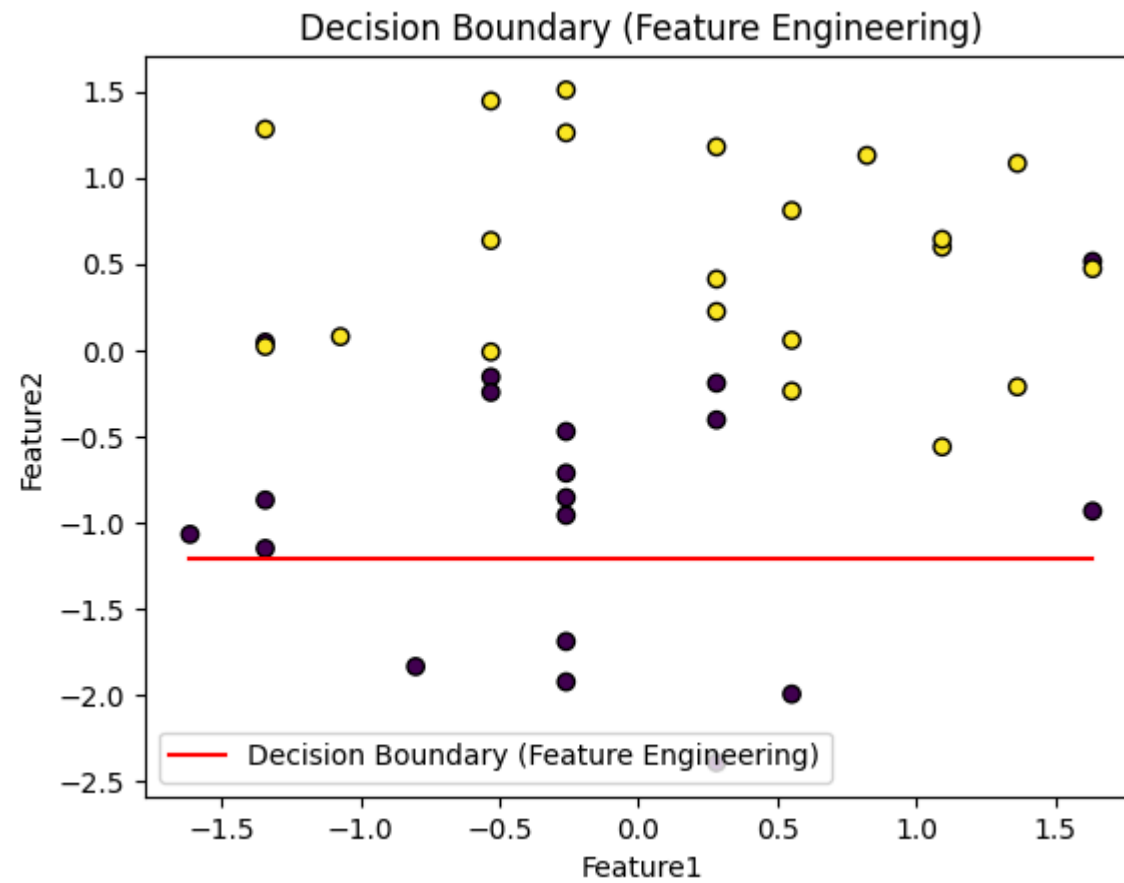
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost Function during Training (Feature Engineering)')
plt.show()

x_values_poly = np.linspace(X_test_poly_scaled[:, 1].min(), X_test_poly_scaled[:, 1].max(), 100)
plt.scatter(X_test_poly_scaled[:, 1], X_test_poly_scaled[:, 2], c=y_test_poly, cmap='viridis', edgecolors='k', marker='o')
y_values_poly = -(trained_theta_poly[0] + trained_theta_poly[1] * x_values_poly) / trained_theta_poly[2]
plt.plot(x_values_poly, y_values_poly, label='Decision Boundary (Feature Engineering)', color='red')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.title('Decision Boundary (Feature Engineering)')
plt.legend()
plt.show()

# Evaluation of Accuracy
predictions_poly = sigmoid(X_test_poly_b @ trained_theta_poly)
predictions_poly = (predictions_poly >= 0.5).astype(int)
accuracy_poly = accuracy_score(y_test_poly, predictions_poly)
print(f"Model Accuracy : {accuracy_poly*100:.2f}%")

```





Model Accuracy : 87.50%

In []: